

1 Supervised Dimensionality Reduction Towards
2 Prediction of Phenotype on
3 Simulated Genotype-Phenotype data

4 Hussain M. Sajwani
5 100045495@ku.ac.ae
6 hussainmsajwani@gmail.com
7 A thesis presented for a Bachelor of Science
8 in Applied Mathematics & Statistics



11 Mathematics Department
12 Khalifa University
 May 2021

13	Contents	
14	1 Abstract	2
15	2 Introduction	2
16	3 Literature review	3
17	4 Methodology	3
18	4.1 PhenotypeSimulator	3
19	4.2 Dimensionality reduction methods	4
20	4.2.1 Neural networks, autoencoders, and supervised	
21	autoencoders	4
22	4.2.2 Principal components analysis (PCA)	8
23	4.2.3 p-value thresholding	8
24	4.2.4 AE weight thresholding	9
25	4.3 Evaluation	9
26	4.3.1 Area under the curve (AUC)	9
27	4.3.2 Precision and Recall	11
28	5 Results	11
29	5.1 Architecture of autoencoder	11
30	5.2 Thresholding	13
31	5.3 Classification	13
32	5.3.1 Classification with PCA reduced dataset	14
33	5.3.2 Classification with SAE reduced dataset	14
34	5.3.3 Classification with thresholded dataset	15
35	6 Discussion	16
36	7 Conclusion	18

1 Abstract

We look at four different methods of dimensionality reduction. Three of which are supervised methods in that they consider a lower-dimensional representation that is more optimal for classification. The methods used are principal components analysis, supervised autoencoder, p -value thresholding, and thresholding based on a supervised autoencoder's weights. We benchmark these methods on simulated datasets where ground truth is known *a priori*. This allows us to fully measure the performance of our models. Furthermore, more we train three different classifiers on the lower-dimensional representation obtained from the four dimensionality reduction methods. We compare the AUC of all three methods to rank the performance of each of the dimensionality reduction methods.

2 Introduction

Height, weight, IQ, eye color, and many others are all traits of any human being. What determines these traits is a mixture of our genetic information and external factors. Besides controlling traits, genetic information and other external factors can also determine whether someone has a disease or not. A natural question arises: given someone's genetic information, can we predict their traits or a disease they might have? To answer this we need access to genetic data. Given how expensive sequencing is and privacy issues, obtaining genetic data is often a hurdle. Even when access to genetic data is obtained, it is often the case that the sample size, n , is much smaller than the dimensionality, d . Classical statistical methods often fail in this high-dimensional setting [1]. Thus if we wanted to use such methods, we will need to find a way of extracting the most important features in our data. In this study, we look at various techniques (supervised autoencoder, weight thresholding, p -value thresholding, PCA) and evaluate their performance on simulated genetic data. In particular, we will extract or construct features that we think are more strongly associated with a particular trait. What we mean here is simulating base positions (locus) that mutate in at least 1% of the population. These loci are called single-nucleotide polymorphisms (SNPs.)

For a given simulated dataset, we also simulate the genotype and the phenotype of n individuals. The genotype of an individual i is a d -dimensional vector representing d SNPs. Let x_{ij} be a random variable denoting the value of the j^{th} SNP of the i^{th} individual. Humans are diploid meaning that at each position of DNA, they have two base pairs, one from each parent. The most common base pair in a population at a particular locus of the DNA is called the major allele. Each x_{ij} takes values in $\{0, 1, 2\}$ representing how much of the major allele of the j^{th} SNP the i^{th} individual has. Let $X \in \mathbb{R}^{n \times d}$ be the matrix containing

the x_{ij} 's mentioned above as its elements. We will call X the *genotype matrix*. The phenotype is then simulated as some combination of genetic effect (genotype), environmental factors, and observational noise [2]. The degree to which each of the three components affects the final phenotype can be varied. The phenotype is an n -dimensional vector $y \in \mathbb{R}^n$ where the i^{th} entry represents a random variable measuring of some trait. The constraint of not having enough data typically means we have $d \gg n$ i.e. we have very few samples compared to how many SNPs we have. This presents a problem as the classical theory of "large n , fixed d " fails to give predictable results. The classical theory breaks down in higher dimensions.

3 Literature review

This idea of using supervised dimensionality reduction is not novel at all. In [3], the authors utilize p -values assessing the goodness of a fit of a logistic model between each SNP and polygenic obesity to judge a SNP's effect in predicting polygenic obesity. The authors then pick a threshold p_0 and use SNPs with $p < p_0$ to predict polygenic obesity using a deep neural network. [3] claims an AUC of 0.99. Another example is the approach in [4], where regions of DNA, called promoters, are passed through a convolutional neural network. The promoters with the highest classification accuracy are then passed into a bigger convolutional neural network. [4] claims an accuracy of 0.769. The authors in [5] propose a novel method of training neural networks in high dimensional regimes. Using some insight from functional analysis, they propose a way of regularizing the weights with the l_p norm by looking at the l_q norm, with $\frac{1}{p} + \frac{1}{q} = 1$ of the neural network's gradients. When this regularization is applied, the authors claim an AUC of 0.926 on a dataset with $n = 85$ and $d = 22,283$.

4 Methodology

To have a good understanding of how models behave, we need to test them out in various settings and scenarios. Testing our models on real data is expensive and time-consuming given how hard it is to get genetic data. We also need ground truth (a priori knowledge) to test our models on. This is why we look at genetic data.

4.1 PhenotypeSimulator

We will use the PhenotypeSimulator (PS) R package [2] to simulate genotype-phenotype relations. PS can either take the genotype as input or simulate its own relatively simple genotype. Since we do not have access to data, we decided to go with the option of simulating the

119 genotypes using PS. The simulated genotypes are very simple. The
 120 genotype matrix X is simulated as

$$\text{vec}(X) \sim \text{Multi}(nd, (p_0, p_1, p_2)).$$

121 Where each p_k represents the allele frequency. e.g. p_0 is the frequency
 122 of there being 0 of the major allele. The phenotype is simulated to be
 123 the output of the genotype and some external factors. Mathematically
 124 speaking, the phenotype is some convex combination of genetic factors
 125 and external factors. In this study, we constrain ourselves to only look-
 126 ing at the genetic factors. In all simulations, we let external factors only
 127 contribute 5% to the simulated phenotype. The genotype contribution
 128 is further broken down into two components. The first controls how
 129 much a select group of SNPs contributes to the phenotype. This select
 130 group of SNPs are called causal SNPs. Typically, this is a very small
 131 amount of SNPs. We will call the number of causal SNPs d_c and their
 132 effect on the final phenotype h_2^s . The second component controls how
 133 much the polygenic effect of all SNPs on the final phenotype. This is
 134 different than the first component in that there are no select few SNPs
 135 that control the phenotype but rather all SNPs contribute. This sec-
 136 ond component will be called h_2^b . Both components are nonnegative real
 137 numbers such that $h_2^s + h_2^b = 1$. We will denote each simulated dataset
 138 by a with 4-tuple

$$(n, d, d_c, h_2^s).$$

139 Unless stated otherwise, we always break down each dataset into a
 140 0.6/0.2/0.2 training, validation, and testing split. This means models
 141 only train on $0.6n$ of the simulated datasets. These will represent the
 142 parameters of interest to us in PS. The phenotype generated is then
 143 some vector $\hat{y} \in \mathbb{R}^n$ which we convert into a binary vector $y \in \{0, 1\}^n$ by
 144 thresholding \hat{y} . An example of this is BMI and obesity. If we are trying
 145 to study obesity, we do not care about the particular value of BMI.
 146 We just care if it is above or below some threshold. A mathematically
 147 inaccurate of how PS works is

$$\text{Phenotype} = 0.95 \times [h_2^s \times \text{Causal SNPs} + h_2^b \times \text{Polygenic effect}] + 0.05 \times \text{External factors}$$

148 4.2 Dimensionality reduction methods

149 In this study, four dimensionality reduction algorithms will be tested.
 150 The first two algorithms will simply extract k features the algorithm
 151 thinks are relevant and discard the rest. The other two will extract k
 152 new features built from the original d features.

153 4.2.1 Neural networks, autoencoders, and supervised autoen- 154 coders

155 A typical binary classifier deep neural network, $\text{DNN} : \mathbb{R}^d \rightarrow [0, 1]$, is
 156 a model made of a sequence of affine transformations that are passed,




Figure 1: A graphical representation of a neural network. Each colored box represents a layer. The connections (black arrows) between the layers are the weights $W^{(l)}$. This image was downloaded from: <https://cs231n.github.io/neural-networks-1/>.

157 elementwise, through some nonlinearities. Let $\{d^{(1)}, \dots, d^{(L)}\} \subset \mathbb{N}$
 158 be a sequence of positive integers. Given an observation $x \in \mathbb{R}^d$, a
 159 DNN apply some affine transformation represented by a matrix $W^{(1)} \in$
 160 $\mathbb{R}^{(d^{(2)}+1) \times (d+1)}$, called the weight matrix, (the +1 is added to represent
 161 the bias term of the affine model. Every observation x is padded with
 162 a constant variable to reflect this bias) and then pass it elementwise
 163 into some nonlinear function, called the activation function, $g^{(1)}$. This
 164 process is repeated L times with the l^{th} time represented with an affine
 165 transformation with matrix $W^{(l)} \in \mathbb{R}^{(d^{(l)}+1) \times (d^{(l-1)}+1)}$ and a nonlinear
 166 function $g^{(l)}$. The l^{th} repetition is called the l^{th} layer of the neural net-
 167 work. $d^{(l)}$ is called the width of the l^{th} layer and the elements of $W^{(l)}$
 168 are called the weights of the layer. The final layer, the L^{th} , has output
 169 size $d^{(L)} = 2$. The output here can be interpreted as two probabilities of
 170 the sample having $y = 0$ or $y = 1$, respectively. Common choices of $g^{(l)}$
 171 are

$$\text{ReLU}(x) = \max\{0, x\}.$$

172 and

$$\sigma(x) = (1 + e^{-x})^{-1}.$$

For a graphical interpretation of a neural network look at figure 1. The choice of L , $d^{(l)}$, and $g^{(l)}$ for any $l = 1, \dots, L$ is made by the designer of a neural network. These parameters which are up to the designer are called the hyperparameters of the neural network. The choice of $W^{(l)}$ for all $l = 1, \dots, L$ though is determined by an optimization algorithm. For a set of weight matrices $W = (W^{(1)}, \dots, W^{(L)})$ we evaluate the performance of the neural network by comparing its prediction $\hat{y} = \text{DNN}(x; W)$ with the actual value of y . Remember that \hat{y} is a two dimensional vector. Its second element \hat{y}_1 is the probability the neural network "thinks" that $y = 1$. In this study we use cross-entropy [6] as our objective function to evaluate the performance of a DNN.




Figure 2: An example of an autoencoder with two layers in the encoder and decoder parts. The red units are the bottleneck. This image was downloaded from <https://www.compthree.com/blog/autoencoder/>

Cross-entropy is defined as

$$\begin{aligned} J(W) &= -\mathbb{E}_{x,y \sim p_{\text{observed}}} \log \text{DNN}(x; W)_1 \\ &= -\log \hat{y}_1. \end{aligned}$$

173 The goal now will be to find W such that $J(W)$ is minimized. First, we
 174 initialize all the weights of all of the neural network randomly. Then
 175 we pass an observation x_i through the network and observe some \hat{y}_i .
 176 With this, we can evaluate $J(W)$ to compare y_i and \hat{y}_i . We can use
 177 this information to judge how good the weights W are. The weights
 178 are then updated to reflect the information just attained. This process
 179 is repeated iteratively until a minimum is reached. We use the Adam
 180 optimizer [7] to find the optimal weights W . The parameters used are
 181 $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$, and $\epsilon = 10^{-7}$.

182 A logistic regression model can be modeled as a one-layer neural
 183 network with $g^{(1)}(x) = (1 + e^{-x})^{-1}$, and $W \in \mathbb{R}^{2 \times (d+1)}$.

184 Autoencoders are a special type of deep neural network where the
 185 final layer is not some binary variable but rather a d -dimensional vector.
 186 Given an input to the neural network $x \in \mathbb{R}^d$, the goal of an autoen-
 187 coder is to first compress x into some smaller vector $x_k \in \mathbb{R}^k$ with $k < d$
 188 and then using x_k to reproduce, to its best ability, the input vector x .
 189 Then let the first $\lfloor L/2 \rfloor$ layers have widths $d^{(1)} > d^{(2)} > \dots > d^{(\lfloor L/2 \rfloor)}$
 190 with $k := d^{(\lfloor L/2 \rfloor)}$. This part of the autoencoder is called the en-
 191 coder. The other layers are to be mirrored around the middle layer.
 192 i.e. $d^{(1)} = d^{(L)}$, $d^{(2)} = d^{(L-1)}$, \dots , and so on. This does not mean these
 193 layers have the same weights. We are only imposing that the widths
 194 of the layers be tied together. Inspired by PCA, there has been some
 195 work in the literature exploring weight-tied autoencoders. See [8] for
 196 a discussion on the matter. This part of the autoencoder is called the
 197 decoder. The middle layer of the autoencoder will be called the latent
 198 space. The final layer then produces a d -dimensional vector \hat{x} which we

interpret as the reconstruction of x . See 2 for a graphical depiction of an autoencoder. Cross-entropy does not work well here. So instead we use mean squared error (MSE) as an objective function to evaluate the performance of the autoencoder. For a vector x that is approximated with \hat{x} by the autoencoder, MSE is defined as $\text{MSE}(W) = (x - \hat{x})^2$. We use Adam [7] to optimize this as well. The compressed representation then comes from taking the output of the middle layer. Notice that these features are extracted features; they might not have a biological interpretation. This is one downside to dimensionality reduction methods that extract features rather than select them. Also, notice that nowhere in the above description do we use the information from the phenotype y . We are simply looking at a lower-dimensional representation of the genotype without taking into account the response of the phenotype. This could be an issue as the autoencoder might not learn to emphasize the presence of a causal SNP. For this reason, we use supervised autoencoders.

We combine the two tasks of classification and reconstruction into a neural network structure called supervised autoencoders [9]. Let

$$\begin{aligned}\text{enc} : \mathbb{R}^d &\longrightarrow \mathbb{R}^k \\ \text{dec} : \mathbb{R}^k &\longrightarrow \mathbb{R}^d \\ \text{clf} : \mathbb{R}^k &\longrightarrow [0, 1].\end{aligned}$$

These functions represent the encoder part of an autoencoder, the decoder part, and a classifier that takes in a k dimensional representation of a d -dimensional vector. Assume that all of these functions are parts of a neural network (each represents a sequence of affine transformations passed through some nonlinear functions.) An autoencoder, AE is then formed by composition

$$AE = \text{dec} \circ \text{enc}.$$

And the supervised part, S , of a supervised autoencoder is formed as

$$S = \text{clf} \circ \text{enc}.$$

We train both of these two parts concurrently. The objective function that we use will be some convex combination of the two previously defined objective functions. As always let W and let $\rho \in [0, 1]$ be the weights of the supervised neural network autoencoder (all three parts of it) then the objective function is defined as

$$J(W) = \rho J_{AE}(W) + (1 - \rho) J_S(W).$$

Where J_{AE} is the objective function of reconstruction and J_S is the objective function of the classifier. We will call the variable ρ the reconstruction weight. It determines how much emphasis the learning

algorithm should put on the reconstruction objective. $1 - \rho$ is the classification weight. It determines the emphasis on the classification objective. This approach to learning is called multitask learning.

As is the case with a classifier deep neural network, the choice of depth and widths is up to the designer to determine. Here we also add the choice of what classifier to use. We look at various choices for these hyperparameters. Unless stated otherwise, assume that `clf` is a logistic regression model.

4.2.2 Principal components analysis (PCA)

Principal components analysis (PCA) is the second reduction method we consider that considers complex features as its extracted features. PCA is an unsupervised method. PCA works by finding a set of orthogonal basis for \mathbb{R}^d such that the variance of the data along each of the basis vectors is maximized. One can think of this as taking a straight line parallel to one of the orthogonal vectors and projecting all of the dataset on it. The variance we try to maximize is the variance of the projected dataset onto the straight line. The theory here is that more variance implies more interesting information. The basis vectors are then sorted by how much variance the dataset has in a vector's direction and we take the top k vectors. Finally, we project the data onto the linear subspace formed by the top k vectors. The projected data is the reduced representation of the original dataset X .

The set of orthogonal basis is found by studying the covariance matrix of X ,

$$\frac{1}{n}XX^T$$

(without loss of generality, assume the variables have mean zero.) Note that this representation of the covariance matrix is made with respect to the canonical basis of \mathbb{R}^d . To maximize variance means to maximize the diagonal of the covariance matrix. This means that we need to diagonalize the matrix to find our sought-after representation. This can be done since covariance matrices are, by definition, symmetric and hence are always diagonalizable and have orthogonal eigenvectors. Finally, the basis we want is the set of eigenvectors of the covariance matrix.

It can be shown that training a two layers autoencoder with linear activation on both layers (i.e. $g^{(1)}(x) = g^{(2)}(x) = x$) and $d^{(2)} = k$ is equivalent to finding the subspace spanned by the top k eigenvectors found above [10].

4.2.3 p-value thresholding

We will look at two thresholding methods. The first one is called p -value thresholding. It is a commonly used method in the field of bioinformatics to determine the strength of the association between a certain SNP

and a certain phenotype. The method fits d logistic regression models between each of the d SNPs and the phenotype y . Then a p -value is calculated to measure the strength of the association between the SNP and the phenotype. That p -value is then used to quantify the strength of the association between a particular SNP and the phenotype. Then we select k SNPs with the k highest value of $-\log_{10} p$. This is called thresholding because typically, one would select some value of $p_0 \in (0, 1)$, called the threshold, and simply select all SNPs with p -values such that $-\log_{10} p \geq -\log_{10} p_0$. However, since we are interested in comparing this method to other methods, we choose a top- k approach.

4.2.4 AE weight thresholding

Another approach we use to select features is based on the autoencoder. After training a supervised autoencoder, we look at the weights of the first layer $W^{(1)} \in \mathbb{R}^{(d^{(2)}+1) \times (d+1)}$ and form a vector,

$$a = \sum_{j=1}^{d^{(2)}+1} \left(W_{:j}^{(1)} \right)^2$$

for the i^{th} SNP. The theory here is that if a SNP is highly associated with the phenotype, then its weights on the first layer should be higher in magnitude. In other words, a_i should be higher. From this, we pick the top k SNPs corresponding to the highest k values of a .

4.3 Evaluation

Since our overarching goal is to predict a phenotype, we will use the performance of classifiers on reduced data to measure the efficacy of our reduction techniques in extracting relevant features. We will use three different classifiers: a support vector machine [11] with a radial basis function kernel, a logistic regression model, and a deep neural network.

The architecture of the DNN used is inspired by the architecture used by [3]. It is 7 layers deep including the input and output layers. The input layer has width $d^{(1)} = k$ (the size of reduced dimensionality.) The rest of the layers are of width 50, 50, 40, 30, 20, 2. To reduce overfitting, we use dropout on the second and third layers. Dropout randomly selects a proportion of the entries of the vector and sets them equal to zero during training [12]. The proportion in this study is 0.5. We also use l_1 and l_2 regularization on all layers. The parameters used for all layers are 3×10^{-6} , and 6.5×10^{-5} .

4.3.1 Area under the curve (AUC)

To measure the performance of the classifiers, we will use a metric called area under the curve (AUC.) The curve in question here is the receiver

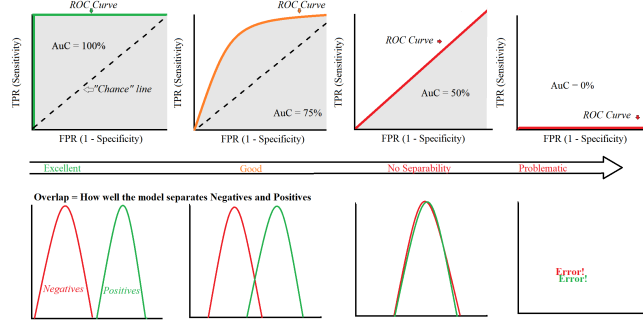


Figure 3: Demonstration of ROC curves. The leftmost panel represents the perfect scenario for a classifier in which it has learnt a perfect separation to the observation based on their class. Given an observation, a classifier with $AUC=100\%$ can always predict which class the observation belongs to. The second panel shows a more attainable ROC curve. The third panel shows a case where the classifier has not learnt anything and leaves classification up to chance. The last panel shows a ROC curve of a model that does worse than chance. It will always give the wrong label. The picture was downloaded from <https://www.datasciencecentral.com/profiles/blogs/roc-curve-explained-in-one-picture>

operating characteristic curve (ROC.) The ROC curve can be thought of as a parametric curve with respect to a thresholding parameter $t \in [0, 1]$. The three methods we test output a number which we interpret as a probability of an observation falling into one of two classes. For any binary classification task, we set some threshold t on the probabilities by which we determine what to label an observation. For certain cases, we require very high certainty so we only label an observation if our models give it a very high probability of being that label so we set a relatively high threshold t . For any choice of threshold, we calculate two parameters. The true positive rate ($TPR(t)$) and the false positive rate ($FPR(t)$). Given a labeled dataset, $TPR(t)$ determines what proportion of observations with a positive label ($y = 1$ in our case) were assigned to have a positive label by the model under consideration and a threshold t . The $FPR(t)$ is the proportion of observations with a negative label ($y = 0$ in our case) were assigned to have a positive label by the model and a threshold t . The ROC curve then is given by

$$\{(FPR(t), TPR(t)) \mid t \in [0, 1].\}$$

A perfect classifier will assign all positive observations a positive label and not assign any negative observations a positive label. This means that the ROC curve will hug the top left corner of the unit square. If we choose our classifier in such a way that it labels each observation

based on the flip of a fair coin, then the ROC curve will be the diagonal line $\text{TPR} = \text{FPR}$. See figure 3 for an illustration. This motivates the definition of AUC. The AUC is defined to be the area under the curve of the ROC curve and above the FPR axis. For a perfect classifier where the curve hugs the top left corner of the unit square, the area under the curve is 1. The fair coin classifier will have AUC 0.5.

We will use AUC to measure the performance of the classifiers.

4.3.2 Precision and Recall

To understand how well our thresholding is doing, we will use two quantities called precision and recall that give us an idea of how well our thresholding is separating causal SNPs from non-causal ones. Precision and recall are typically used to assess classifiers. We treat thresholding as a classifier that assigns a positive label to quantities above (or below) the threshold and negative labels to quantities below (or above) the threshold. As was the case with the true positive rate and the false positive rate, precision and recall are functions of the thresholding parameter t . Precision is defined as

$$\text{Prec}(t) = \frac{\text{TPR}(t)}{\text{TPR}(t) + \text{FPR}(t)}.$$

It is interpreted as the proportion of all observations that have been assigned a positive label that are positive. Recall is defined as

$$\text{Rec}(t) = \frac{\text{TPR}(t)}{\text{TPR}(t) + \text{FNR}(t)}.$$

Where $\text{FNR}(t)$ is the rate of positive observations that are assigned a negative label. Recall is interpreted as the proportion of all positive observations that have been labeled correctly.

5 Results

Computations were carried on a Linux machine running Ubuntu 18.04.5 with 64GB RAM, 2 Nvidia Tesla K20c GPUs, and 12 Xeon E5-2603 cores. Miniconda3 Python 3.7.9 is used for all computations. Neural network based models have been implemented in TensorFlow 2.2.0 [13]. Other models used have been implemented using scikit-learn 0.23.2 [14].

5.1 Architecture of autoencoder

On the left of figure 4 we see a histogram of how much each possible occurrence of the major allele is present at any SNP. We see that most loci (location on DNA) have 0 occurrences of the major allele. Fewer loci have 1 occurrence, and even fewer loci have 2 occurrences. We simulate

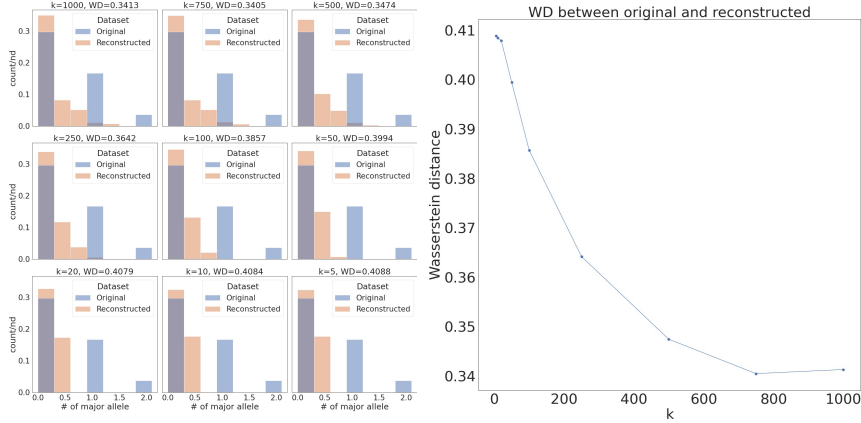


Figure 4: *Left:* Each panel has a histogram of the original genotype matrix X and a histogram of \hat{X} , the autoencoder’s attempt at reconstructing X . The genotype is simulated with parameters $(600, 1000, 10, 0.5)$. The autoencoder is two layers deep with $\rho = 0.8$ with $k \in \{5, 10, 20, 50, 100, 250, 500, 750, 1000\}$ *Right:* Wasserstein distance from original histogram to reconstructed ones from autoencoders with size k .

a dataset with parameters $(600, 1000, 10, 0.5)$ and train 9 supervised autoencoders with $\rho = 0.8$ and values of k ranging from 1000 and 5. It can be seen from the histograms that as k decreases (i.e. the bottleneck tightens), the autoencoder learns to just focus on reconstructing as many zeros as possible.

To quantify the observed effect, we utilize the Wasserstein distance [15]. The Wasserstein distance is a way of measuring distances between two probability distributions. It can be shown that the Wasserstein measure defines a metric on the set of all Borel measures on some metric space. This allows us to measure the distance between the original histogram of the dataset and the autoencoder’s attempt at reconstructing it. The rightmost panel of figure 4 shows the Wasserstein distance between the original histogram and the reconstruction. This confirms what we see in the middle panel. As the bottleneck size k decreases, the reconstruction becomes harder to learn.

To further understand the learnt supervised autoencoder, SAE, we look at the extracted features. Specifically, we look at the covariances of the extracted features and how that covariance varies with k and ρ . For this, we train a few ReLU activated SAEs with $\rho \in \{0.2, 0.6, 1\}$ and different k values on a dataset. The results shown in figure 5 show the correlation matrices of the extracted features. It can be seen that as




Figure 5: Correlation matrices of extracted features of ReLU activated SAE. Each panel represents the results of a SAE trained with $\rho = 0.2, 0.6$, and 1 , respectively. Each panel contains nine subpanels which show the correlation matrices for values of $k \in \{5, 10, 20, 50, 100, 250, 500, 750, 1000\}$. The leftmost, $\rho = 1$, panel confirms theory that an AE will learn a diagonal correlation matrix, like PCA. The other two SAE with $\rho \neq 1$ do not learn uncorrelated features. Note that if an entry on the diagonal is zero, then the SAE learnt a constant feature.

382 ρ approaches 1 i.e. the SAE approaches an AE, the correlation matrix
383 tends to become more diagonal. If the SAE has learnt to extract a
384 constant feature, then the variance of that feature will be zero and
385 hence correlation will be undefined. When this happens we let the
386 correlation be 0.

387 5.2 Thresholding

388 To compare the two thresholding methods we look at how well they
389 capture causal SNPs in various settings of h_2^s . To do this we utilize the
390 precision and recall quantities defined above. We simulate 125 datasets
391 with parameters $(600, 1000, 10, h_2^s)$ for $h_2^s \in \{0.05, 0.25, 0.5, 0.75, 1\}$ and
392 perform both types of thresholding on each of the 25 datasets. The SAE
393 used has $k = 300$ and $\rho = 0.3$. The k used for thresholding is taken to
394 be $k \in \{5, 25, 100, 300\}$. Figure 7 shows an example of manhattan plots
395 where the x axis is all of the SNPs and the y represents how strongly
396 associated the SNP is to the phenotype. Figure 6 shows violin plots
397 of precision and recall over all 25 datasets for each value of h_2^s . The
398 straight lines mean that there was no variance in the quantity being
399 measured.

400 5.3 Classification

401 In the following sections, we use the same 125 simulated datasets from
402 the thresholding section. Here we benchmark the dimensionality reduction
403 methods on classifiers. We use the same 125 simulated datasets in
404 section 5.2 and apply the four reduction methods discussed. We reduce

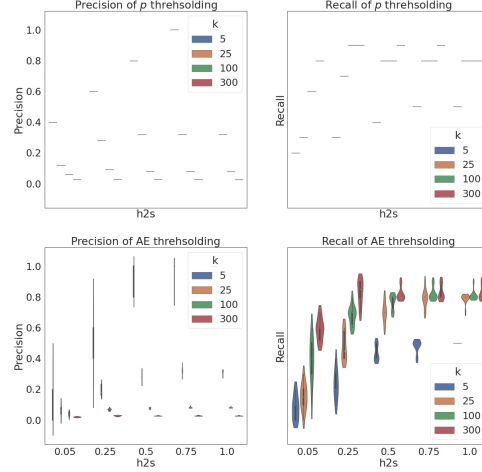


Figure 6: Violin plots of precision and recall of both thresholding methods over all 25 datasets for each value of h_2^s and k . In both precision plots, precision goes up as h_2^s increases and goes down as k increases. Recall increase with both k and h_2^s . *Top left*: precision of p value thresholding. In all circumstances, p thresholding has zero variance. *Top right*: recall of p thresholding. This also has zero variance. *Bottom left*: precision of AE weights thresholding. High variance for lower values of k and decreases with h_2^s . *Bottom right*: Recall of AE thresholding. Low variance on lower values of k .

the dataset to k features with $k \in \{5, 25, 100, 300\}$. Then we train three classifier models on each of the reduced representations for every k . The classifier models to be used are the same ones mentioned in 4.3

5.3.1 Classification with PCA reduced dataset

Figure 8 shows the classification results of the three models trained on the projected dataset with different values of h_2^s . The SVM and logistic model seem to outperform the deep neural network. The AUC does not exceed 0.9 in all of the models. With respect to each model, the performance of the model does not change with h_2^s . Increasing the value of k improves performance.

5.3.2 Classification with SAE reduced dataset

Figure 9 shows the results of passing through the representation a SAE learns in its middle layer to a classification algorithm. SAE seems to

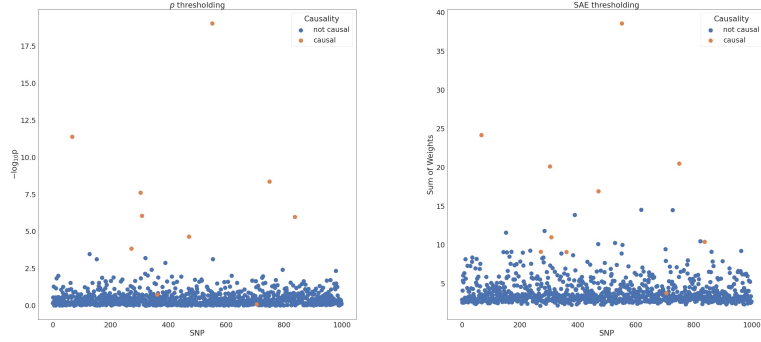


Figure 7: Example thresholding on a single dataset with $h_2^s = 0.75$ and both thresholding methods.




Figure 8: Classification results of three models trained after reducing dimensionality using PCA. In all three models, PCA has done a bad job of extracting information that is useful for classification. The AUC barely goes above 0.8 in some cases.

be more sensitive to h_2^s than PCA. Each color, on average rises up as h_2^s is increased. Furthermore, in lower k 's, SAE's results seem to outperform PCA's. Even though the supervised branch is made of a logistic regression model, the separately trained logistic model on the learnt data representation seems to have a different outcome. The separately trained logistic model has high variance in lower k 's.

5.3.3 Classification with thresholded dataset

The results of thresholding are shown in figure 10. The left column shows the results of p -value thresholding while the right column shows the results of the SAE thresholded one. p -value thresholding blows both in bias (AUC) and variance. The performance of SAE thresholding falls quite drastically as k increases. The neural network seems to be struggling with high variance with both thresholding methods.




Figure 9: Classification results of three models trained after reducing dimensionality using SAE with $\rho = 0.3$. All methods seem to be doing better than PCA. *Top left:* the performance of the supervised part of the supervised autoencoder. *Top right:* results of an SVM trained on the latent space of the autoencoder

6 Discussion

The first thing we want to determine in an autoencoder is how small can we make the latent space while still preserving the integrity of the reconstructed dataset. We will first need to understand how the simulated datasets look like. Ideally, we would want our autoencoder to recreate this structure when attempting to reconstruct the dataset. To do this we look at the results in 4 and try to find a compromise between k and preserving the structure of the dataset. Choosing $k \geq 500$ does not yield a lot of dimensionality reduction. As can be seen 4, the SAE is having a hard time reconstructing the rare case of having two major alleles for smaller values of k . We see that for $k \geq 250$ the SAE is doing a decent job of reconstructing the original structure. Hence, we choose

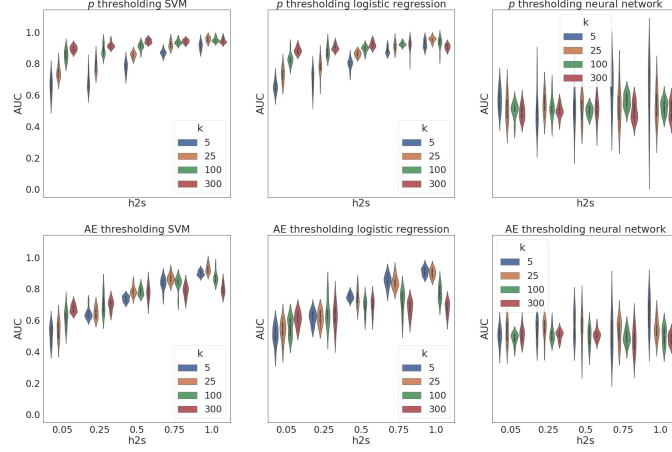


Figure 10: Classification results of three models trained after reducing dimensionality using thresholding.

a value of $k \in [250, 500]$.

We know from theory that linearly activated autoencoders have the same results as PCA. The correlation matrix of PCA’s learnt representation is a diagonal matrix. Even though our supervised autoencoders use a ReLU activation, from 5 they still tend to a diagonal matrix as ρ goes to 1. For values of ρ near 1, the supervised autoencoder does not pay much attention to the supervised part and hence approaches a vanilla autoencoder with no supervised part. SAEs with lesser values of ρ however, seem to learn more correlated features. Unlike PCA or a vanilla autoencoder, a SAE does not just care about representing as much variance as possible. This can be seen in 5 where SAEs with low ρ learnt correlated features.

This emphasis on classification is reflected by how the improvement in AUC of classifiers that learn on SAE data rather than PCA data. SAE classifiers have better AUC on average than PCA classifiers as is seen in figures 8 and 9. However, the problem of high variance in performance persists in SAEs. While PCA might suffer from not being able to extract relevant features for classification, it is still more consistent than a SAE. One way to reduce this high variance is to employ a lot of regularization on the weights of the neural network. This can be done with techniques like the ones mentioned in [5]. A future piece of work is to incorporate extreme regularization methods on SAEs.

Figures 7 and 6 give us more insight on how the AE weights thresholding method compares to the method of p -value thresholding. The

key insight here is that AE thresholding has quite a lot more variance than p thresholding. The AE method will often label noncausal SNPs as causal. This can be particularly seen in figure 7 where p thresholding does a better job of separating causal SNPs from noncausal ones. This is understandable given how many more parameters an AE has compared to a simple logistic model.

This can also be seen in the AUC results in figure 10. p thresholding yields a better selection of SNPs and hence classifiers do a better job. Even when it is easiest to separate causal SNPs from noncausal SNPs, when h_2^s is near 1, the AUC drops drastically as k increases. This can be understood by looking at 7 where adding more SNPs will add more noise (noncausal SNPs) to the model. Keep in mind that our datasets have no linkage disequilibrium. So no correlations between SNPs are made. We hypothesize that the gap between p thresholding and AE thresholding will decrease if we add more linkage disequilibrium. This is because p thresholding only looks at single SNPs but not the interaction of multiple SNPs.

7 Conclusion

Overall the two methods which select features have outperformed the other two methods which extract features. The p thresholding method has proven to be better than thresholding with the weights of a SAE. SAEs outperformed PCA in extracting features relevant for classification. SAEs were relatively cheap to train (200 MB of VRAM.) Hence if the goal of dimensionality reduction is to reduce the computational cost, then SAEs are a better alternative than PCA.

References

- [1] M. Wainwright, *High-Dimensional Statistics: A Non-Asymptotic Viewpoint*. Cambridge University Press, 02 2019.
- [2] H. V. Meyer and E. Birney, “PhenotypeSimulator: A comprehensive framework for simulating multi-trait, multi-locus genotype to phenotype relationships,” *Bioinformatics*, vol. 34, pp. 2951–2956, 03 2018.
- [3] C. A. C. Montañez, P. Fergus, A. C. Montañez, A. Hussain, D. Al-Jumeily, and C. Chalmers, “Deep learning classification of polygenic obesity using genome wide association study snps,” *CoRR*, vol. abs/1804.03198, 2018.
- [4] B. Yin, M. Balvert, R. A. A. van der Spek, B. E. Dutilh, S. Bohté, J. Veldink, and A. Schönhuth, “Using the structure of genome data in the design of deep neural networks for predicting amyotrophic lateral sclerosis from genotype,” *Bioinformatics*, vol. 35, pp. i538–i547, 07 2019.
- [5] B. Liu, Y. Wei, Y. Zhang, and Q. Yang, “Deep neural networks for high dimension, low sample size data,” in *Proceedings of the 26th International Joint Conference on Artificial Intelligence, IJCAI’17*, p. 2287–2293, AAAI Press, 2017.
- [6] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [7] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” 2017.
- [8] P. Vincent, H. Larochelle, I. Lajoie, Y. Bengio, and P.-A. Manzagol, “Stacked denoising autoencoders: Learning useful representations in a deep network with a local denoising criterion,” *J. Mach. Learn. Res.*, vol. 11, p. 3371–3408, Dec. 2010.
- [9] L. Le, A. Patterson, and M. White, “Supervised autoencoders: Improving generalization performance with unsupervised regularizers,” in *Advances in Neural Information Processing Systems* (S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, eds.), vol. 31, Curran Associates, Inc., 2018.
- [10] P. Baldi and K. Hornik, “Neural networks and principal component analysis: Learning from examples without local minima,” *Neural Networks*, vol. 2, no. 1, pp. 53–58, 1989.
- [11] G. James, D. Witten, T. Hastie, and R. Tibshirani, *An Introduction to Statistical Learning: with Applications in R*. Springer Texts in Statistics, Springer New York, 2013.

- 532 [12] N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and
533 R. Salakhutdinov, “Dropout: A simple way to prevent neu-
534 ral networks from overfitting,” *J. Mach. Learn. Res.*, vol. 15,
535 p. 1929–1958, Jan. 2014.
- 536 [13] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro,
537 G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat,
538 I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Joze-
539 fowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga,
540 S. Moore, D. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner,
541 I. Sutskever, K. Talwar, P. Tucker, V. Vanhoucke, V. Vasude-
542 van, F. Viégas, O. Vinyals, P. Warden, M. Wattenberg, M. Wicke,
543 Y. Yu, and X. Zheng, “TensorFlow: Large-scale machine learning
544 on heterogeneous systems,” 2015. Software available from tensor-
545 flow.org.
- 546 [14] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion,
547 O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg,
548 J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot,
549 and E. Duchesnay, “Scikit-learn: Machine learning in Python,”
550 *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830,
551 2011.
- 552 [15] P. Clement and W. Desch, “An elementary proof of the triangle
553 inequality for the wasserstein metric,” *Proceedings of the American*
554 *Mathematical Society*, vol. 136, no. 1, pp. 333–339, 2008.